

AI is the new interface layer. SuperAlign secures it.

Discover, score and govern every AI interaction — at the source · Datasheet

Your network tools watch the wire and the cloud — downstream of where AI actually thinks. SuperAlign puts an agentic sensor, **Surface**, on the endpoint, browser and network path to see and govern every agent, model and MCP server the moment it acts — including the third-party and local AI nothing else can see.

#1

reason CISO security budgets are rising: securing AI

63%

of AI usage never touches the network

15_{min}

time to value — complete visibility across the fleet

1M+

AI assets pre-scored in Atlas, 8 categories

Budget finding: Felicis AI-security survey, 2025 (45 enterprise security leaders, 76% CISO/CSO-level). Coverage from Atlas / AIRiskDB; other figures illustrative.

THE SHIFT THAT BROKE THE THREAT MODEL

PRE-AI ERA

User

 →

Application

Employees used public SaaS and private apps directly through the UI. The **application was the control point**, and security wrapped around it.

CASB

SSO

IAM

DLP

POST-AI ERA

User

 →

AI

 →

Everything

AI is the **new interface layer** — the OS, the browser, the front door to everything. The app no longer mediates access; the AI does. **Security lost visibility.**

No control point

Blind spot

This is a fundamental shift in **where the control point lives** — and in the threat model around it.

01

One prompt, many systems

A single request can touch Salesforce, internal databases and third-party APIs in one motion.

02

One agent, many tools

Agents invoke tools and skills on the user's behalf, often with no human in the loop.

03

One MCP server, many capabilities

A single connector can expose far more surface area than the org realizes it has granted.

04

The user is the new source of truth

The application no longer mediates access — the user, through the AI, does.

WHY THE EXISTING STACK CAN'T SEE IT

Coverage depends on where you sit in the AI data path. Everything downstream of the device reads yesterday's logs to stop today's action. Surface sits at the earliest point — the endpoint itself.

COVERAGE	SURFACE · ENDPOINT	GATEWAY / SASE	SDK IN YOUR CODE	CASB / DLP
Local models, MCP servers, on-device agents	● Complete	● Never sees it	● Out of scope	● Never sees it
Browser & IDE AI assistants	● Complete	● URL filtering only	● Out of scope	● Sanctioned apps
Cloud AI egress	● With full context	● In transit, no context	● First-party only	● Known apps only
Off-network & remote devices	● Always on	● Only if routed	● n/a	● Only via proxy
Agent intent & the action taken	● Classified on device	● None	● App-level only	● None
Enforcement point	● At source, pre-action	● Downstream, in transit	● Your own apps	● At the proxy

Category comparison — positions in the AI-security landscape, not specific vendors.

THE PLATFORM · ONE SENSOR, THREE PRODUCTS, ONE CONSOLE

One agentic sensor — **SuperAlign Surface** — plus a browser extension and network telemetry feeds three products, all driven from a single console. SuperAlign governs AI **interactions**: what your people and their agents do with AI. It does not touch your first-party apps or non-AI traffic.



AGENTIC ENDPOINT SENSOR

SuperAlign Surface

One lightweight agentic sensor on every endpoint, browser and network path — the vantage point that sees every AI interaction, online or off. It discovers the assets, carries the telemetry, and enforces the verdict on the device, before the action completes.



RISK INTELLIGENCE · AIRISKDB

SuperAlign Atlas

Risk intelligence behind every verdict — so the moment an asset appears, its risk is known.

✓ 1M+ assets pre-scored across 8 categories

✓ Capability, Exposure & Governance dimensions

✓ 57 signal definitions; CVSS-aligned 0–10



DETECTION & RESPONSE · AIDR

SuperAlign Sentinel

Purpose-built small models run locally, scoring intent and risk in milliseconds.

✓ Intent classification at the edge, not domains

✓ Enforce allow / warn / block before the action

✓ 87 scanner signals; private, works offline



SYSTEM OF RECORD · AI-IEM SOC

SuperAlign Signal

The AI-native system of record. A SIEM indexes logs; Signal captures cognition.

✓ Prompt-level observability of every AI event

✓ Replayable, span-by-span agent traces

✓ Immutable audit, streamed into your SOC



CONTROL PLANE

SuperAlign Console

One pane of glass — policy, investigation and response across every sensor.

✓ Set guardrails once; enforce on every endpoint

✓ Pivot from a risk score to a trace to the device

✓ Board-ready reporting, mapped to frameworks

EVERY ASSET SCORED ON THREE RISK DIMENSIONS

Capability

what the asset can do & access

- Web crawling, code execution, file-system access, autonomous actions, tool & API reach

Exposure

footprint across systems & users

- Install reach, public inference endpoints, always-on, cross-platform availability

Governance

legal, compliance & policy

- Legal & compliance status, enterprise-account usage, risk-framework alignment, policy violations, user privilege

- macOS
- Windows
- Linux



15 min

first AI inventory

- Managed browser extension
- Deploys through existing MDM
- No re-architecture or rerouting

Mapped to EU AI Act · ISO 42001 · NIST AI RMF · MITRE ATLAS · OWASP LLM, Agent, MCP & Skill Top 10

Framework-mapped

See your own AI inventory and gather telemetry within 15 minutes.

cal.com/appidi